

Introduction to Econometrics

Moinul Islam

9/16/2022

Table of contents

Preface	4
1 Nature and scope of Econometrics	5
2 Introductory Econometrics	6
2.1 Install packages	6
2.2 Useful dataset	6
2.3 Simple regression model	6
2.3.1 Example 1: A log wage equation	6
2.4 Multiple regression analysis	10
2.4.1 Example 2: Hourly wage equation	10
2.5 Multiple regression analysis: inference	12
2.5.1 Example 3: Hourly Wage Equation	12
2.5.2 Example 4: Effect of Job Training on Firm Scrap Rates	16
2.6 Chapter 5: Multiple Regression Analysis: OLS Asymptotics	20
2.6.1 Example: Housing Prices and Distance From an Incinerator	20
3 Statistical Inference	25
4 The Simple Regression Model	26
5 Multiple Regression Analysis: Estimation	27
6 Multiple Regression Analysis: Inference	28
7 Multiple Regression Analysis: Further Topics	29
8 Asymptotic properties of OLS	30
9 Heteroskedasticity	31
10 Time Series data Modelling Dynamic Processes	32
11 Autocorrelation	33
12 Endogeneity – omitted variables, measurement error and simultaneity	34

13 Instrumental variable estimation and two stage least squares	35
14 Logit and Probit Models for Binary Response	36
15 Non Stationary Time Series	37
16 Simultaneous Equations Models	38
References	39

Preface

1 Nature and scope of Econometrics

2 Introductory Econometrics

2.1 Install packages

```
library(wooldridge)
library(readr)
library(stargazer)
library(kableExtra)
library(quantmod)
library(xts)
```

2.2 Useful dataset

```
earnings <- read.csv("data/earnings.csv")
gpa1 <- read.csv("data/gpa1.csv")
hprice1 <- read.csv("data/hprice1.csv")
hprice2 <- read.csv("data/hprice2.csv")
hprice3 <- read.csv("data/hprice3.csv")
jtrain <- read.csv("data/jtrain.csv")
nyse <- read.csv("data/nyse.csv")
phillips <- read.csv("data/phillips.csv")
rdchem <- read.csv("data/rdchem.csv")
traffic1 <- read.csv("data/traffic1.csv")
wage1 <- read.csv("data/wage1.csv")
```

2.3 Simple regression model

2.3.1 Example 1: A log wage equation

- Load the `wage1` data and check out the documentation.

educ: years of education


wage: average hourly earnings

lwage: log of the average hourly earnings

- First, make a scatter-plot of the two variables and look for possible patterns in the relationship between them.

```
plot(y = wage1$wage, x = wage1$educ, col = "black", pch = 21, bg = "orange",  
     cex=1.25, xaxt="n", frame = FALSE, main = "Wages vs. Education, 1976",  
     xlab = "years of education", ylab = "Hourly wages")  
axis(side = 1, at = c(0,6,12,18))  
rug(wage1$wage, side=2, col="black")
```

Wages vs. Education, 1976



1. It appears that *on average*, more years of education, leads to higher wages.
 2. The example in the text is interested in the *return to another year of education*, or what the *percentage* change in wages one might expect for each additional year of education. To do so, one must use the *log(wage)*. This has already been computed in the data set and is defined as *lwage*.
- Build a linear model to estimate the relationship between the *log of wage* (*lwage*) and *education* (*educ*).

$$\widehat{\log(wage)} = \beta_0 + \beta_1 educ$$

```
log_wage_model <- lm(lwage ~ educ, data = wage1)
```

- Print the summary of the results.

```
summary(log_wage_model)
```

Call:

```
lm(formula = lwage ~ educ, data = wage1)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.21158	-0.36393	-0.07263	0.29712	1.52339

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.583773	0.097336	5.998	3.74e-09 ***
educ	0.082744	0.007567	10.935	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4801 on 524 degrees of freedom

Multiple R-squared: 0.1858, Adjusted R-squared: 0.1843

F-statistic: 119.6 on 1 and 524 DF, p-value: < 2.2e-16

- Use the stargazer package to make beautiful table

```
stargazer(type = "html", log_wage_model, single.row = TRUE, header = FALSE, digits = 3)
```

Dependent variable:

lwage

educ

0.083*** (0.008)

Constant

0.584*** (0.097)

Observations

526

R2

0.186

Adjusted R2

0.184

Residual Std. Error

0.480 (df = 524)

F Statistic


119.582*** (df = 1; 524)

Note:

$p < 0.1$; $p < 0.05$; $p < 0.01$

- Plot the $\log(\text{wage})$ vs educ. The blue line represents the least squares fit.

```
plot(y = wage1$lwage, x = wage1$educ, main = "A Log Wage Equation",  
     col = "orange", pch = 21, bg = "black", cex=1.25,  
     xlab = "years of education", ylab = "log of average hourly wages",  
     xaxt="n", frame = FALSE)  
axis(side = 1, at = c(0,6,12,18))  
abline(log_wage_model, col = "black", lwd=2)  
rug(wage1$lwage, side=2, col="black")
```



2.4 Multiple regression analysis

2.4.1 Example 2: Hourly wage equation

Check the documentation for variable information

lwage: log of the average hourly earnings


educ: years of education

exper: years of potential experience

tenure: years with current employer

- Plot the variables against *lwage* and compare their distributions and slope (β) of the simple regression lines

```
par(mfrow=c(1,3))
plot(y = wage1$lwage, x = wage1$educ, col="orange", xaxt="n", frame = FALSE, main = "years of education")
mtext(side=2, line=2.5, "Hourly wages", cex=1.25)
axis(side = 1, at = c(0,6,12,18))
abline(lm(lwage ~ educ, data=wage1), col = "black", lwd=2)
plot(y = wage1$lwage, x = wage1$exper, col="orange", xaxt="n", frame = FALSE, main = "years of experience")
axis(side = 1, at = c(0,12.5,25,37.5,50))
abline(lm(lwage ~ exper, data=wage1), col = "black", lwd=2)
plot(y = wage1$lwage, x = wage1$tenure, col="orange", xaxt="n", frame = FALSE, main = "years with employer")
axis(side = 1, at = c(0,11,22,33,44))
abline(lm(lwage ~ tenure, data=wage1), col = "black", lwd=2)
```



- Estimate the model regressing *educ*, *exper*, and *tenure* against $\log(\text{wage})$.

$$\log(\widehat{wage}) = \beta_0 + \beta_1 educ + \beta_3 exper + \beta_4 tenure$$

```
hourly_wage_model <- lm(lwage ~ educ + exper + tenure, data = wage1)
```

- Print the estimated model coefficients:


```
coefficients(hourly_wage_model)
```

```
kable(coefficients(hourly_wage_model), digits=4, col.names = "Coefficients", align = 'l')
```

	Coefficients
(Intercept)	0.2844
educ	0.0920
exper	0.0041
tenure	0.0221

- Plot the coefficients, representing percentage impact of each variable on $\log(wage)$ for a quick comparison.

```
barplot(sort(100*hourly_wage_model$coefficients[-1]), horiz=TRUE, las=1,
        ylab = " ", main = "Coefficients of Hourly Wage Equation")
```



2.5 Multiple regression analysis: inference

2.5.1 Example 3: Hourly Wage Equation

Using the same model estimated in **example 3**, examine and compare the standard errors associated with each coefficient. Like the textbook, these are contained in parenthesis next to each associated coefficient.

```
summary(hourly_wage_model)
```

Call:

```
lm(formula = lwage ~ educ + exper + tenure, data = wage1)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.05802	-0.29645	-0.03265	0.28788	1.42809

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.284360	0.104190	2.729	0.00656 **
educ	0.092029	0.007330	12.555	< 2e-16 ***
exper	0.004121	0.001723	2.391	0.01714 *
tenure	0.022067	0.003094	7.133	3.29e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4409 on 522 degrees of freedom

Multiple R-squared: 0.316, Adjusted R-squared: 0.3121

F-statistic: 80.39 on 3 and 522 DF, p-value: < 2.2e-16

```
stargazer(type = "html", hourly_wage_model, single.row = TRUE, header = FALSE, digits=5)
```

Dependent variable:

lwage

educ

0.09203*** (0.00733)

exper

0.00412** (0.00172)

tenure

0.02207*** (0.00309)

Constant

0.28436*** (0.10419)

Observations

526

R2

0.31601

Adjusted R2

0.31208

Residual Std. Error

0.44086 (df = 522)

F Statistic

80.39092*** (df = 3; 522)

Note:

$p < 0.1$; **$p < 0.05$** ; $p < 0.01$

For the years of experience variable, or **exper**, use coefficient and Standard Error to compute the t statistic:

$$t_{\text{exper}} = \frac{0.004121}{0.001723} = 2.391$$

Fortunately, R includes t statistics in the **summary** of model diagnostics.

```
summary(hourly_wage_model)$coefficients
```




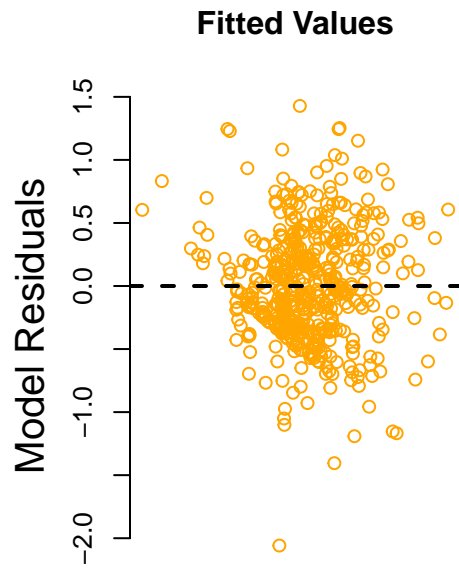
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.284359541	0.104190379	2.729230	6.562466e-03
educ	0.092028988	0.007329923	12.555246	8.824197e-32
exper	0.004121109	0.001723277	2.391437	1.713562e-02
tenure	0.022067218	0.003093649	7.133070	3.294407e-12

```
kable(summary(hourly_wage_model)$coefficients, align="l", digits=3)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.284	0.104	2.729	0.007
educ	0.092	0.007	12.555	0.000
exper	0.004	0.002	2.391	0.017
tenure	0.022	0.003	7.133	0.000


- lets plot this results

```
par(mfrow=c(2,2))
plot(y = hourly_wage_model$residuals, x = hourly_wage_model$fitted.values , col="orange",
     frame = FALSE, main = "Fitted Values", xlab = "", ylab = "")
mtext(side=2, line=2.5, "Model Residuals", cex=1.25)
abline(0, 0, col = "black", lty=2, lwd=2)
plot(y = hourly_wage_model$residuals, x = wage1$educ, col="green", xaxt="n",
     frame = FALSE, main = "years of education", xlab = "", ylab = "")
axis(side = 1, at = c(0,6,12,18))
abline(0, 0, col = "black", lty=2, lwd=2)
plot(y = hourly_wage_model$residuals, x = wage1$exper, col="gray", xaxt="n",
     frame = FALSE, main = "years of experience", xlab = "", ylab = "")
mtext(side=2, line=2.5, "Model Residuals", cex=1.25)
axis(side = 1, at = c(0,12.5,25,37.5,50))
abline(0, 0, col = "black", lty=2, lwd=2)
plot(y = hourly_wage_model$residuals, x = wage1$tenure, col="blue", xaxt="n",
     frame = FALSE, main = "years with employer", xlab = "", ylab = "")
axis(side = 1, at = c(0,11,22,33,44))
abline(0, 0, col = "black", lty=2, lwd=2)
```



- Plot the t statistics for a visual comparison:

```
barplot(sort(summary(hourly_wage_model)$coefficients[-1, "t value"]), horiz=TRUE, las=1,
        ylab = " ", main = "t statistics of Hourly Wage Equation")
```



2.5.2 Example 4: Effect of Job Training on Firm Scrap Rates

- Load the `jtrain` data set. (From H. Holzer, R. Block, M. Cheatham, and J. Knott (1993), *Are Training Subsidies Effective? The Michigan Experience*, Industrial and Labor Relations Review 46, 625-636. The authors kindly provided the data.)

year : 1987, 1988, or 1989

union : =1 if unionized

lscrap : Log(scrap rate per 100 items)

hrsemp : (total hours training) / (total employees trained)

lsales : Log(annual sales, \$)

lemploy : Log(umber of employees at plant)

- First, use the `subset` function and it's argument by the same name to return observations which occurred in **1987** and are not **union**.
- At the same time, use the `select` argument to return only the variables of interest for this problem.

```
jtrain_subset <- subset(jtrain, subset = (year == 1987 & union == 0), select = c(year, uni
```


- Next, test for missing values. One can “eyeball” these with R Studio’s **View** function, but a more precise approach combines the `sum` and `is.na` functions to return the total number of observations equal to NA.

```
sum(is.na(jtrain_subset))
```

[1] 156


- While R’s `lm` function will automatically remove missing NA values, eliminating these manually will produce more clearly proportioned graphs for exploratory analysis. Call the `na.omit` function to remove all missing values and assign the new `data.frame` object the name `jtrain_clean`.

```
jtrain_clean <- na.omit(jtrain_subset)
```


- We use `jtrain_clean` to plot the variables of interest against `lscrap`. Visually observe the respective distributions for each variable, and compare the slope (β) of the simple regression lines.

```
par(mfrow=c(1,3))
point_size <- 1.60
plot(y = jtrain_clean$lscrap, x = jtrain_clean$hrsemp, frame = FALSE,
main = "Total (hours/employees) trained", ylab = "", xlab="", pch = 21, bg = "lightgrey",
mtext(side=2, line=2, "Log(scrap rate)", cex=1.25)
abline(lm(lscrap ~ hrsemp, data=jtrain_clean), col = "black", lwd=2)
plot(y = jtrain_clean$lscrap, x = jtrain_clean$lsales, frame = FALSE, main = "Log(annual s
abline(lm(lscrap ~ lsales, data=jtrain_clean), col = "black", lwd=2)
plot(y = jtrain_clean$lscrap, x = jtrain_clean$lemploy, frame = FALSE, main = "Log(# emplo
abline(lm(lscrap ~ lemploy, data=jtrain_clean), col = "black", lwd=2)
```


Total (hours/employees) trai



Log(annual sales \$)



Log(# employees at plant)



- Now create the linear model regressing `hrsemp`(total hours training/total employees trained), `lsales`(log of annual sales), and `lemploy`(the log of the number of the employees), against `lscrap`(the log of the scrape rate).

$$lscrap = \alpha + \beta_1 hrsemp + \beta_2 lsales + \beta_3 lemploy$$

```
linear_model <- lm(lscrap ~ hrsemp + lsales + lemploy, data = jtrain_clean)
```

- Finally, print the complete summary diagnostics of the model.

```
summary(linear_model)
```

Call:

```
lm(formula = lscrap ~ hrsemp + lsales + lemploy, data = jtrain_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.6301	-0.7523	-0.4016	0.8697	2.8273

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.45837	5.68677	2.191	0.0380 *
hrsemp	-0.02927	0.02280	-1.283	0.2111
lsales	-0.96203	0.45252	-2.126	0.0436 *
lemploy	0.76147	0.40743	1.869	0.0734 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.376 on 25 degrees of freedom

Multiple R-squared: 0.2624, Adjusted R-squared: 0.1739

F-statistic: 2.965 on 3 and 25 DF, p-value: 0.05134

- Use `stargazer` to create representative table

```
stargazer(type = "html", linear_model, single.row = TRUE, header = FALSE, digits=5)
```

Dependent variable:

lscrap

hrsemp

-0.02927 (0.02280)

lsales

-0.96203** (0.45252)

lemploy

0.76147* (0.40743)

Constant

12.45837** (5.68677)

Observations

29

R2

0.26243

Adjusted R2

0.17392

Residual Std. Error

1.37604 (df = 25)

F Statistic


2.96504* (df = 3; 25)

Note:

$p < 0.1$; $p < 0.05$; $p < 0.01$

```
coefficient <- coef(linear_model)[-1]
confidence <- confint(linear_model, level = 0.95)[-1,]
graph <- drop(barplot(coefficient, ylim = range(c(confidence)),
                     main = "Coefficients & 95% C.I. of variables on Firm Scrap Rates"))
arrows(graph, coefficient, graph, confidence[,1], angle=90, length=0.55, col="black", lwd=2)
arrows(graph, coefficient, graph, confidence[,2], angle=90, length=0.55, col="black", lwd=2)
```

Coefficients & 95% C.I. of variables on Firm Scrap Rates



2.6 Chapter 5: Multiple Regression Analysis: OLS Asymptotics

2.6.1 Example: Housing Prices and Distance From an Incinerator

- We will use the `hprice3` data set.


lprice : Log(selling price)

ldist : Log(distance from house to incinerator, feet)

larea : Log(square footage of house)

- Graph the prices of housing against distance from an incinerator:

```
par(mfrow=c(1,2))
plot(y = hprice3$price, x = hprice3$dist, main = " ", xlab = "Distance to Incinerator in f
abline(lm(price ~ dist, data=hprice3), col = "black", lwd=2)
```



- Next, model the $\log(\text{price})$ against the $\log(\text{dist})$ to estimate the percentage relationship between the two.

$$\text{price} = \alpha + \beta_1 \text{dist}$$

```
price_dist_model <- lm(lprice ~ ldist, data = hprice3)
```

- Create another model that controls for “quality” variables, such as square footage **area** per house.

$$\text{price} = \alpha + \beta_1 \text{dist} + \beta_2 \text{area}$$

```
price_area_model <- lm(lprice ~ ldist + larea, data = hprice3)
```

- Compare the coefficients of both models. Notice that adding **area** improves the quality of the model, but also reduces the coefficient size of **dist**.

```
summary(price_dist_model)
```

Call:

```
lm(formula = lprice ~ ldist, data = hprice3)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.22356	-0.28076	-0.05527	0.27992	1.29332

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8.25750	0.47383	17.427	< 2e-16 ***
ldist	0.31722	0.04811	6.594	1.78e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4117 on 319 degrees of freedom

Multiple R-squared: 0.1199, Adjusted R-squared: 0.1172

F-statistic: 43.48 on 1 and 319 DF, p-value: 1.779e-10

```
summary(price_area_model)
```

Call:

```
lm(formula = lprice ~ ldist + larea, data = hprice3)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.23380	-0.18820	-0.01723	0.21751	0.86039

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.49394	0.49065	7.121	7.18e-12 ***
ldist	0.19623	0.03816	5.142	4.77e-07 ***
larea	0.78368	0.05358	14.625	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3188 on 318 degrees of freedom

Multiple R-squared: 0.4738, Adjusted R-squared: 0.4705

F-statistic: 143.2 on 2 and 318 DF, p-value: < 2.2e-16

- Use *stargazer* for better table

```
stargazer(type = "html",price_dist_model, price_area_model, single.row = TRUE, header = F
```

Dependent variable:

lprice

(1)

(2)

ldist

0.31722*** (0.04811)

0.19623*** (0.03816)

larea

0.78368*** (0.05358)

Constant

8.25750*** (0.47383)

3.49394*** (0.49065)

Observations

321

321

R2

0.11994

0.47385

Adjusted R2

0.11718

0.47054

Residual Std. Error

0.41170 (df = 319)

0.31883 (df = 318)

F Statistic

43.47673*** (df = 1; 319)



143.19470*** (df = 2; 318)

Note:

$p < 0.1$; $p < 0.05$; $p < 0.01$

- Graphing illustrates the larger coefficient for area

Log(distance from incinerator) Log(square footage of house)



3 Statistical Inference

4 The Simple Regression Model

5 Multiple Regression Analysis: Estimation

6 Multiple Regression Analysis: Inference

7 Multiple Regression Analysis: Further Topics

8 Asymptotic properties of OLS

9 Heteroskedasticity

10 Time Series data Modelling Dynamic Processes

11 Autocorrelation

12 Endogeneity – omitted variables, measurement error and simultaneity

13 Instrumental variable estimation and two stage least squares

14 Logit and Probit Models for Binary Response

15 Non Stationary Time Series

16 Simultaneous Equations Models

References